

Library Management System Design

1. Functional Requirements

(What the system DOES)

User Authentication:

Users can register, log in, and reset passwords.

Librarians have admin access (add/update/delete books, manage users).

Book Search:

Search books by title, author, ISBN, or genre.

View real-time availability status.

Borrowing/Returning:

Borrow up to 5 books/user.

Automatic due date calculation (14-day loan period).

Return books with overdue fines (20 pesos/day).

Notifications:

Email/SMS reminders 3 days before due dates.

Alerts for overdue books.

Admin Features:

Add/update/remove books from the catalog.

Generate reports (popular books, overdue items).

2. Non-Functional Requirements

(How the system PERFORMS)

Performance:

Search results load in <1.5 seconds (even under peak load).

The system handles 300+ concurrent users with load balancing (NGINX) and caching.

Scalability:

Database supports 100,000+ books and 50,000+ users.

Designed for horizontal scaling (additional servers can be added via PostgreSQL sharding or read replicas).

Implements read/write database separation (e.g., primary-replica database setup) to improve query performance.

Uses caching (Redis/Memcached) for frequently accessed book data to reduce database load.

Security:

Passwords stored using **bcrypt** or **Argon2id** (salted, adaptive hashing).

HTTPS encryption for all data transfers.

Regular security audits.

Availability:

99% uptime (excluding scheduled maintenance).

Usability:

Intuitive UI for users with no technical background.

Reliability:

Daily automated backups.

Transaction rollback on errors (e.g., failed borrow/return).

3. Subsystems

(Modular Components)

Subsystem	Responsibilities	Key Parameters
<i>User Management & Authentication</i>	<ul style="list-style-type: none">- Registration, login, profile updates.- Role-based access (user vs. librarian).- Session management.	<ul style="list-style-type: none">- Max 5 books/user.- Password complexity rules.- JWT token expiration: 24 hours.
<i>Book Inventory</i>	<ul style="list-style-type: none">- Add/update/remove books.- Track availability, ISBN, metadata.	<ul style="list-style-type: none">- Unique ISBN per book.- Categorize by genre.
<i>Borrowing System</i>	<ul style="list-style-type: none">- Process loans/returns.- Calculate due dates and fines.	<ul style="list-style-type: none">- 14-day loan period.- 20 pesos/day overdue fine.
<i>Notification System</i>	<ul style="list-style-type: none">- Send reminders via email/SMS.- Alert librarians about overdue books.	<ul style="list-style-type: none">- 3-day pre-due reminder.- Daily overdue alerts.
<i>Reporting Module</i>	<ul style="list-style-type: none">- Generate usage statistics.- Export lists of overdue books.	<ul style="list-style-type: none">- Weekly/Monthly reports.

4. Key Parameters

Max Books/User: 5 (adjustable by librarians).

Database Scalability:

Initial capacity: 50,000 books, 10,000 users.

Scalable to 500,000 books and 100,000 users using sharding (PostgreSQL Citus) or cloud databases (AWS Aurora).

Performance Considerations:

Caching: Redis for frequent searches (e.g., popular genres).

Indexing: Optimize search with indexes on title, author, ISBN.

Security Parameters:

HTTPS with TLS 1.3.

Regular security audits.

Password policy: Minimum 8 characters with symbols/numbers.

Session timeout: 15 minutes of inactivity.

Database Model Comparison: Relational vs. NoSQL

<i>Criteria</i>	Relational (SQL)	NoSQL
<i>Data Structure</i>	Structured, table-based with fixed schema (rows and columns).	Flexible schema (document, key-value, graph, or wide-column stores).
<i>Scalability</i>	Horizontal (sharding, read replicas) + Vertical.	Horizontal scaling (adding more servers).
<i>Consistency</i>	ACID compliance (Atomicity, Consistency, Isolation, Durability).	BASE model (Basically Available, Soft state, eventually consistent).
<i>Query Language</i>	SQL (powerful for complex joins and transactions).	Varies by type (e.g., MongoDB uses JSON-like queries; limited joins).
<i>Use Cases</i>	Systems requiring complex transactions (e.g., banking, inventory).	High-speed, unstructured data (e.g., social media, IoT, real-time analytics).

<i>Performance</i>	Optimized for read-heavy operations and complex queries.	Faster for write-heavy operations and unstructured data.
<i>Maintenance</i>	Requires schema migrations; strict data integrity.	Schema-less design allows flexibility but risks inconsistent data without governance.
<i>Examples</i>	MySQL, PostgreSQL, SQL Server.	MongoDB (document), Cassandra (wide-column), Redis (key-value).

Trade-offs

Model	Strengths	Weaknesses
<i>Relational</i>	<ul style="list-style-type: none"> - ACID guarantees. - Complex queries with SQL. - Hybrid scaling. 	<ul style="list-style-type: none"> - Schema changes require migrations.
<i>NoSQL</i>	<ul style="list-style-type: none"> - Horizontal scalability. - Flexible schema. - High write performance. 	<ul style="list-style-type: none"> - No atomic transactions. - Manual joins for relationships

Justification for Choosing Relational Database

Scenario: Library management systems require transactional integrity (e.g., borrowing a book update both user loans and book availability) and structured data (fixed relationships between users, books, and loans).

Scalability

Why Relational (PostgreSQL)?

- Horizontal scaling: Achieved via read replicas (for search queries) and sharding (e.g., Citus extension).
- Cloud compatibility: AWS Aurora/Google Cloud SQL automate scaling for 100,000+ books.

- **Vertical scaling:** Upgrading server specs (CPU/RAM) handles initial growth (50k books).

Why Not NoSQL?

- **Eventual consistency:** Risk of stale data (e.g., a book appears available but is already borrowed).
- **Transactional limits:** NoSQL cannot atomically update `User.borrowed_books_count` and `Book.availability_status`.

Maintainability

Why Relational?

- **Structured schema:** Clear relationships (e.g., `User` ↔ `Loan` ↔ `Book` tables).
- **Data governance:** Constraints (e.g., foreign keys) enforce rules like “max 5 books/user”.
- **Schema migrations:** Tools like Liquibase/Flyway automate schema changes (e.g., adding a new column to the `Book` table).
- **JSONB support:** Store semi-structured book metadata (e.g., genres, tags) while retaining ACID.

Why Not NoSQL?

- **Schema-less design:** Risk of inconsistent data formats (e.g., some books lack genre or author fields).
- **Manual joins:** Relationships (e.g., linking loans to users and books) require application-level code, increasing maintenance complexity.
- **No built-in constraints:** Rules like “max 5 books/user” must be enforced in code, increasing error risk.
- **Data cleanup:** Requires manual effort to fix inconsistencies (e.g., orphaned loan records).

Functional Fit

ACID Compliance: Critical for loan transactions (e.g., deducting a book’s availability and updating a user’s loan count in one atomic operation).

Complex Queries: SQL simplifies generating reports (e.g., “Top 10 most borrowed books”) with JOIN and GROUP BY.

Final Recommendation

Chosen Model: Relational Database (PostgreSQL)

Justification:

- Ensures transactional integrity for loans/returns.
- Structured schema aligns with library data relationships.
- Hybrid scaling (horizontal + vertical) meets 500,000-book scalability needs.
- JSONB allows flexibility for semi-structured metadata.

The **Library Management System** was initially designed as a **simple** solution for book borrowing, returning, and searching. However, to ensure **real-world usability and long-term efficiency**, the design was refined with **scalability, performance, and security enhancements**. Instead of just a basic system, it now supports **500,000 books and 100,000 users**, leveraging **PostgreSQL sharding and read replicas** for horizontal scaling. To maintain **fast search speeds (<1.5s)** and **support 300+ concurrent users**, **Redis caching, indexing, and NGINX load balancing** were integrated. Security was also strengthened with **bcrypt/Argon2id** password hashing, **HTTPS encryption, and session management**. With **99% uptime, automated backups, and ACID compliance**, the system is built not just for functionality but for **real-world reliability**, proving that even a "simple" system should be **designed for scalability and maintainability** to meet modern requirements.